


Base exploitation rises with scale; ground-truth training removes it



n=80 held-out prompts, 8 generations. Configs differ across rungs (7B: 200 steps @ $\beta=0.05$), so this is not a controlled scaling study.